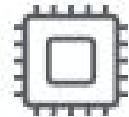


# AWS SERVICES OVERVIEW

# AWS Compute Services



## Compute

EC2

Elastic Beanstalk

ECR

ECS

EKS

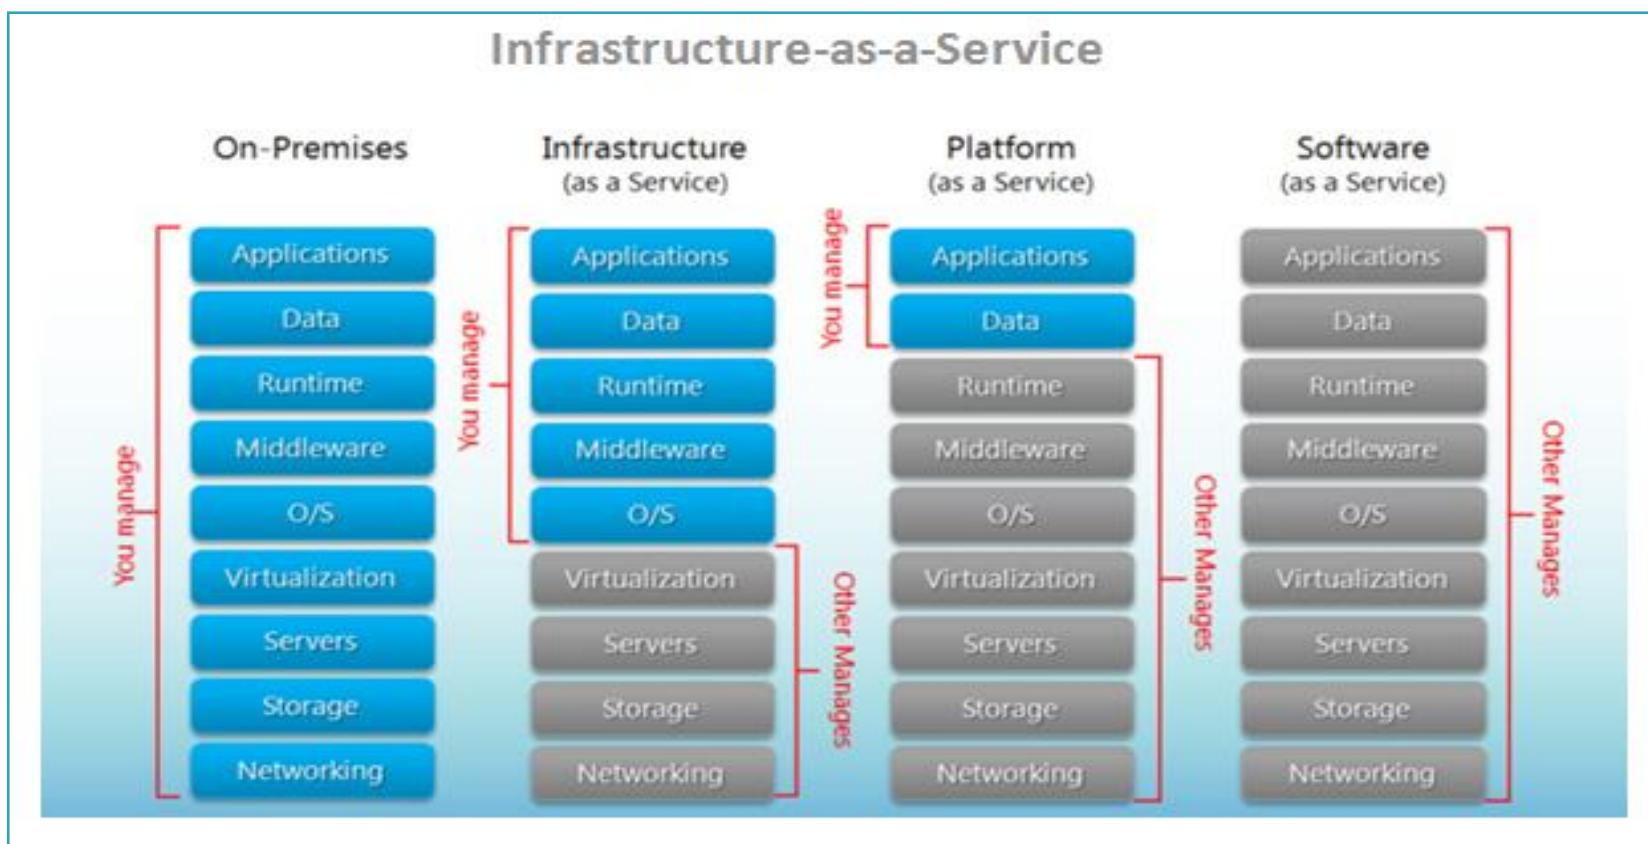
Lambda

Batch

Lightsail

# EC2 (Elastic Compute Cloud)

Amazon Elastic Compute Cloud (Amazon EC2) provides scalable computing capacity in the Amazon Web Services (AWS) cloud



# EC2 Overview



Amazon  
EC2

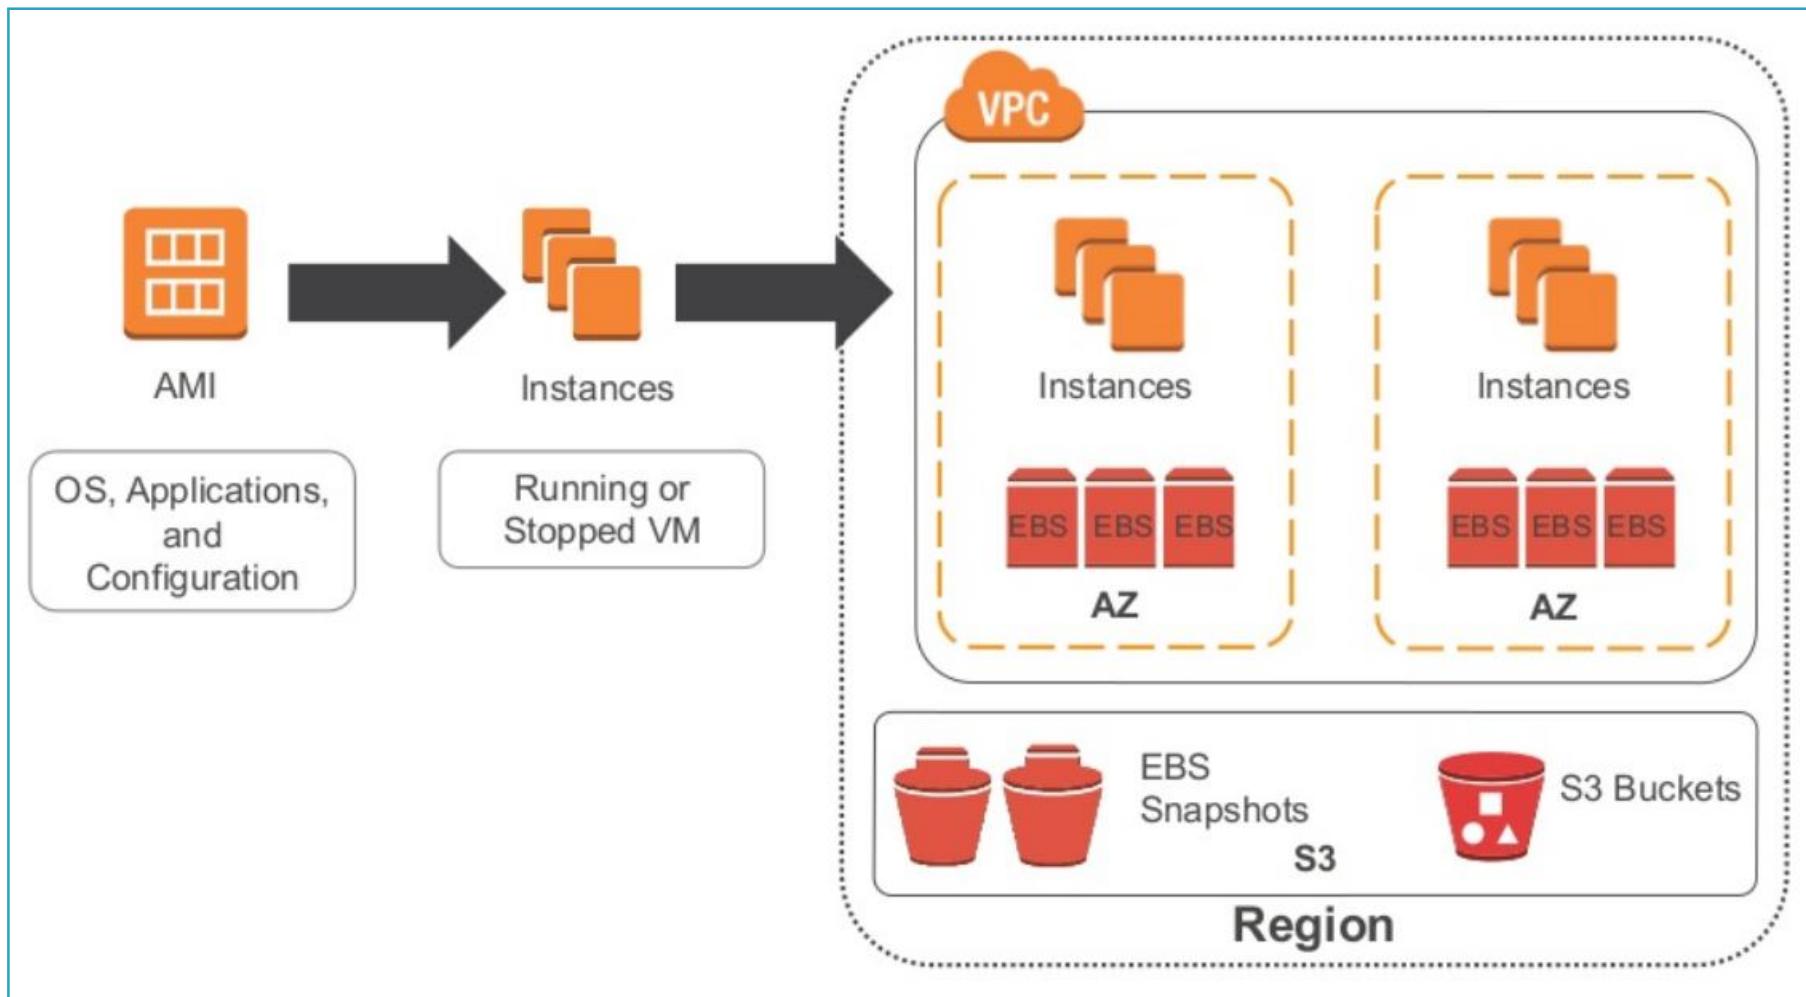
- **Resizable** compute capacity
- Complete control of your computing resources
- **Reduced time required** to obtain and boot new server instances

- **Scale capacity** as your computing requirements change
- Pay only for capacity that you actually use
- Choose **Linux** or **Windows**
- Deploy across **AWS Regions** and **Availability Zones** for reliability
- Use **tags** to help manage your Amazon EC2 resources

# EC2 Concepts

- **Instance** - Virtual Computing environments
- **AMI** - Preconfigured template of instance
- **Instance Types** - Various configs of CPU, memory, storage and network
- **Key Pair** – Secure login to instance
- **Instance Store Volumes** – Temporary Data Storage
- **EBS (Elastic Block Store)** – Persistent Storage Volume
- **Availability Zones** – Physical locations of resources
- **Security Groups** – Firewall to protect resource access based on rules
- **Placement Groups** - Determines how instances are placed on underlying hardware
- **Tags** – Metadata about the instance

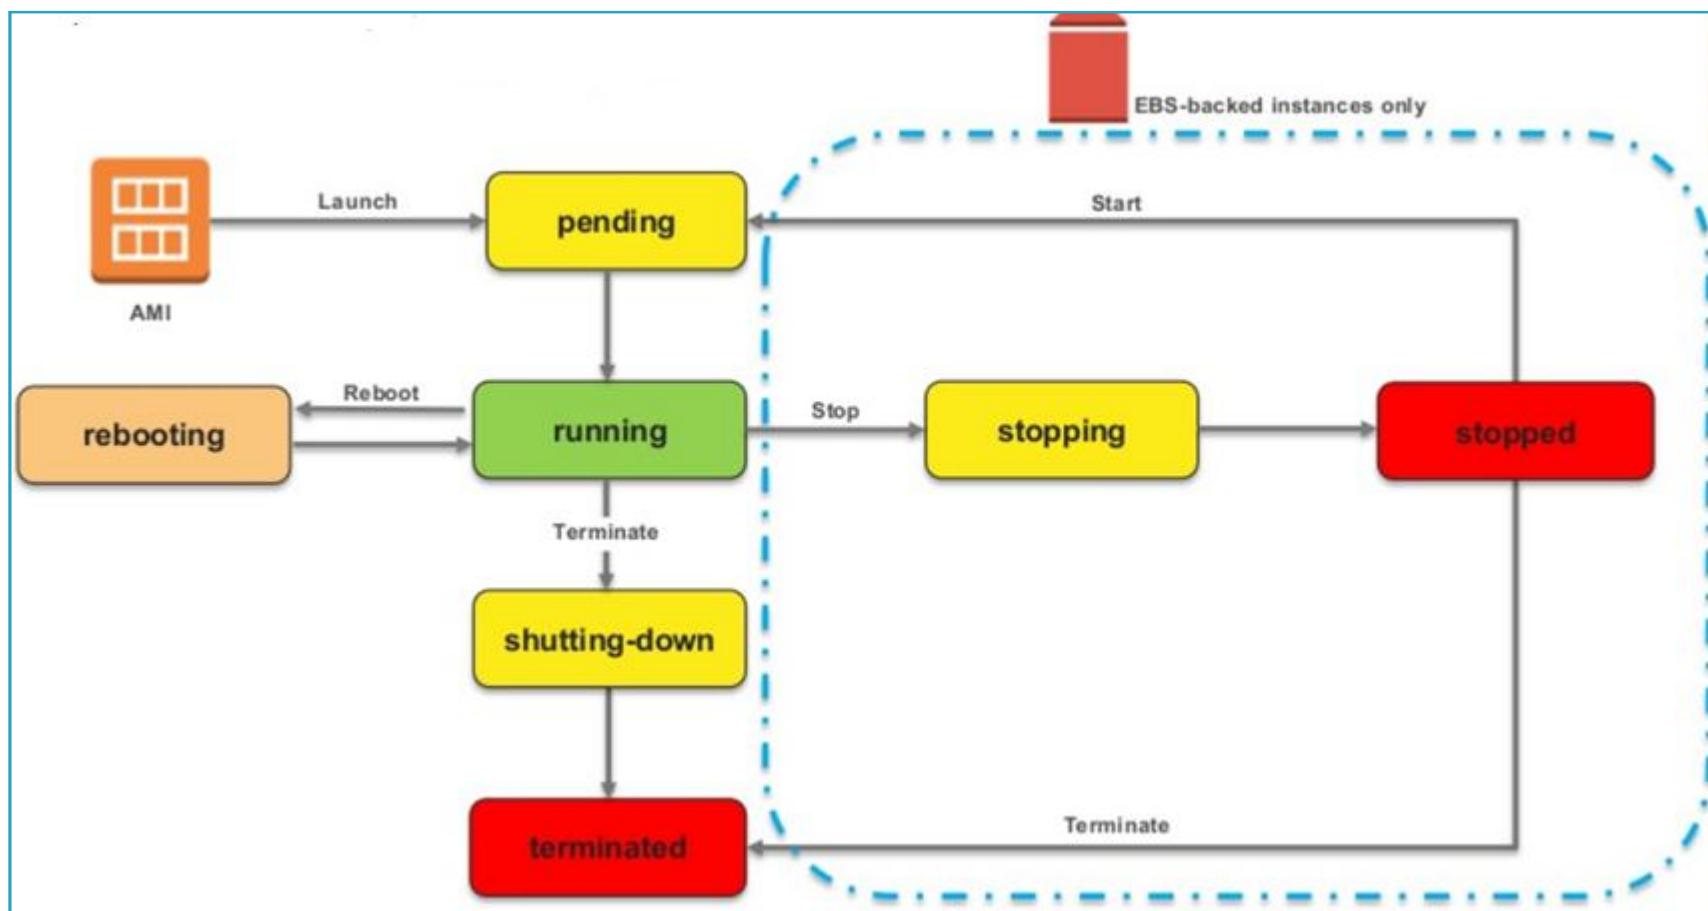
# EC2 Instances



# EC2 Instance Types

| Instance Family              | Some Use Cases   |
|------------------------------|--|
| General purpose (t2, m4, m3) | <ul style="list-style-type: none"><li>• Low-traffic websites and web applications</li><li>• Small databases and mid-size databases</li></ul> |
| Compute-optimized (c4, c3)   | <ul style="list-style-type: none"><li>• High performance front-end fleets</li><li>• Video-encoding</li></ul>                                 |
| Memory-optimized (r3)        | <ul style="list-style-type: none"><li>• High performance databases</li><li>• Distributed memory caches</li></ul>                             |
| Storage-optimized (i2, d2)   | <ul style="list-style-type: none"><li>• Data warehousing</li><li>• Log or data-processing applications</li></ul>                             |
| GPU instances (g2)           | <ul style="list-style-type: none"><li>• 3D application streaming</li><li>• Machine learning</li></ul>  |

# Instance Lifecycle



# Instance Purchasing Options

## On-Demand Instances

Pay by the hour.

## Reserved Instances

Purchase, at a significant discount, instances that are always available

1-year to 3-year terms.

## Scheduled Instances

Purchase instances that are always available on the specified recurring schedule, for a one-year term.

## Spot Instances

Bid on unused instances, which can run as long as they are available and your bid is above the Spot price.

## Dedicated Instances

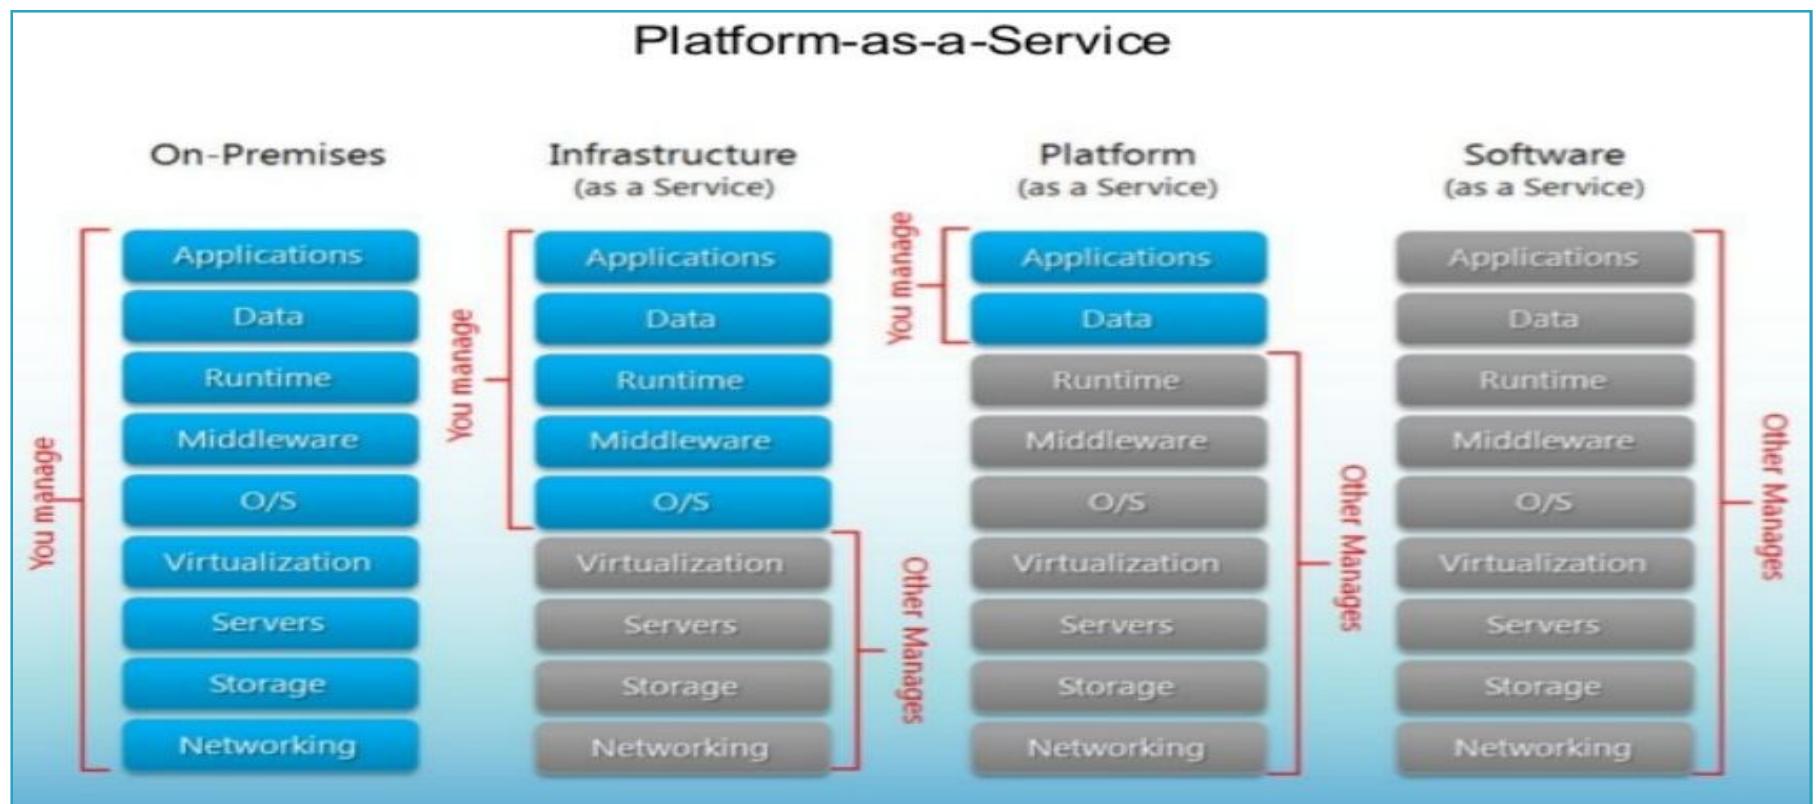
Pay, by the hour, for instances that run on single-tenant hardware.

## Dedicated Hosts

Pay for a physical host that is fully dedicated to running your instances.

# Elastic Beanstalk

**ElasticBeanstalk** enables to quickly deploy and manage applications in the AWS Cloud without worrying about the infrastructure that runs those applications. Reduces management complexity without restricting choice or control.



# Elastic Beanstalk concepts

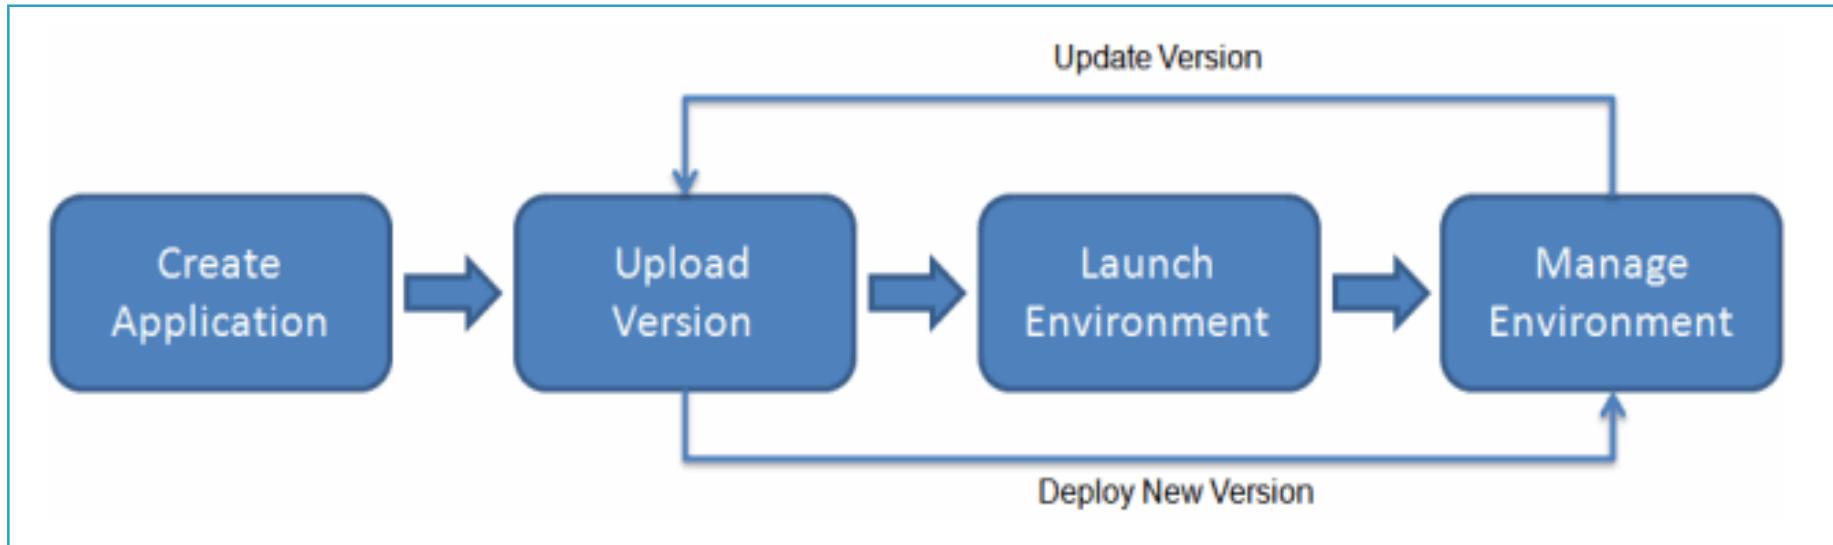
- **Application** - Logical collection of Beanstalk components (env, version, config)
- **Application Version** - Specific, labeled iteration of deployable code
- **Environment** - Provisions the resources needed to run the application version
- **Environment Tier** – Web server or Worker environment
- **Environment Config** – Collection of parameters or settings to define behavior
- **Configuration Template** – Blueprint of environment configuration

# Elastic Beanstalk Launguage Support

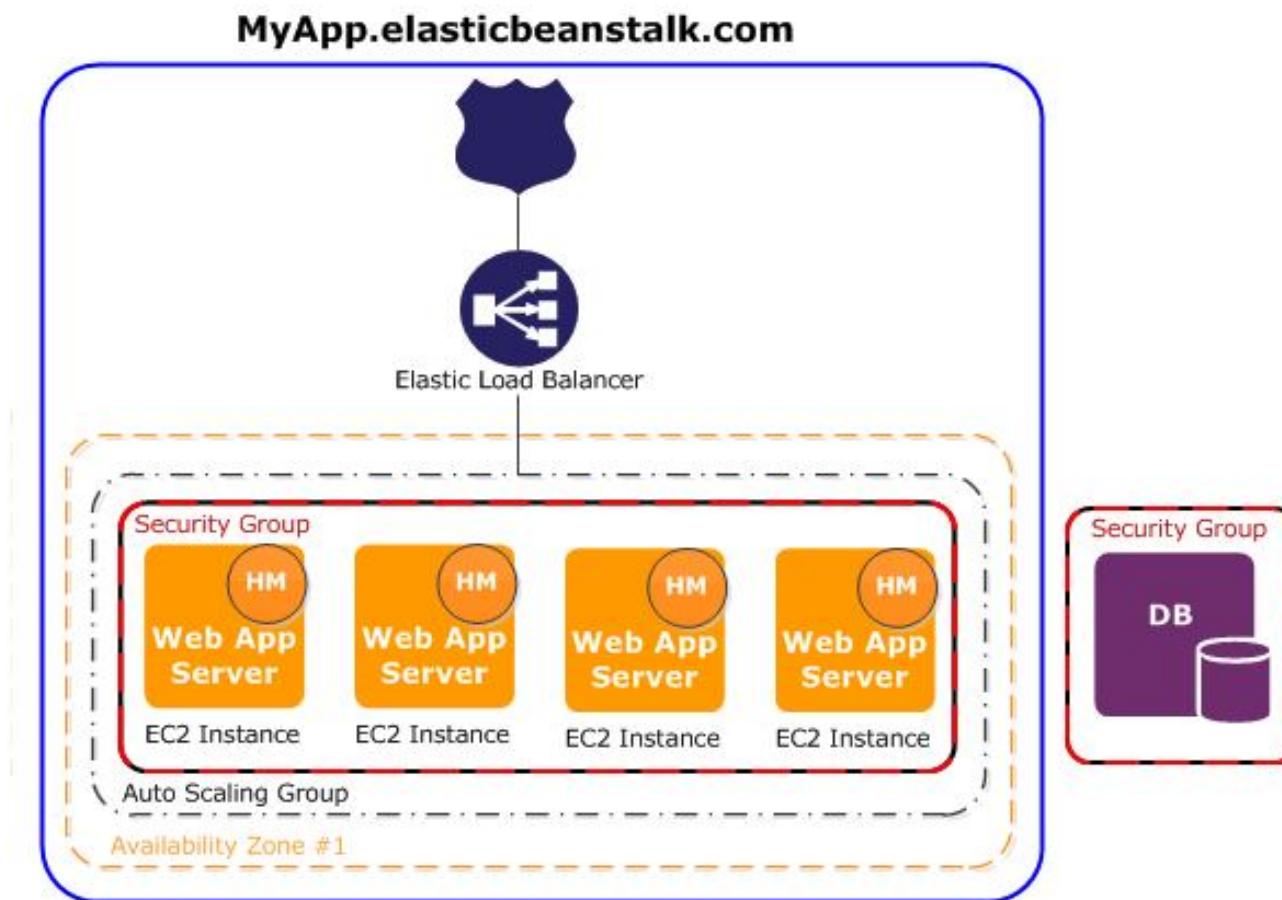


**Deployment Artifacts:** zip, war, docker

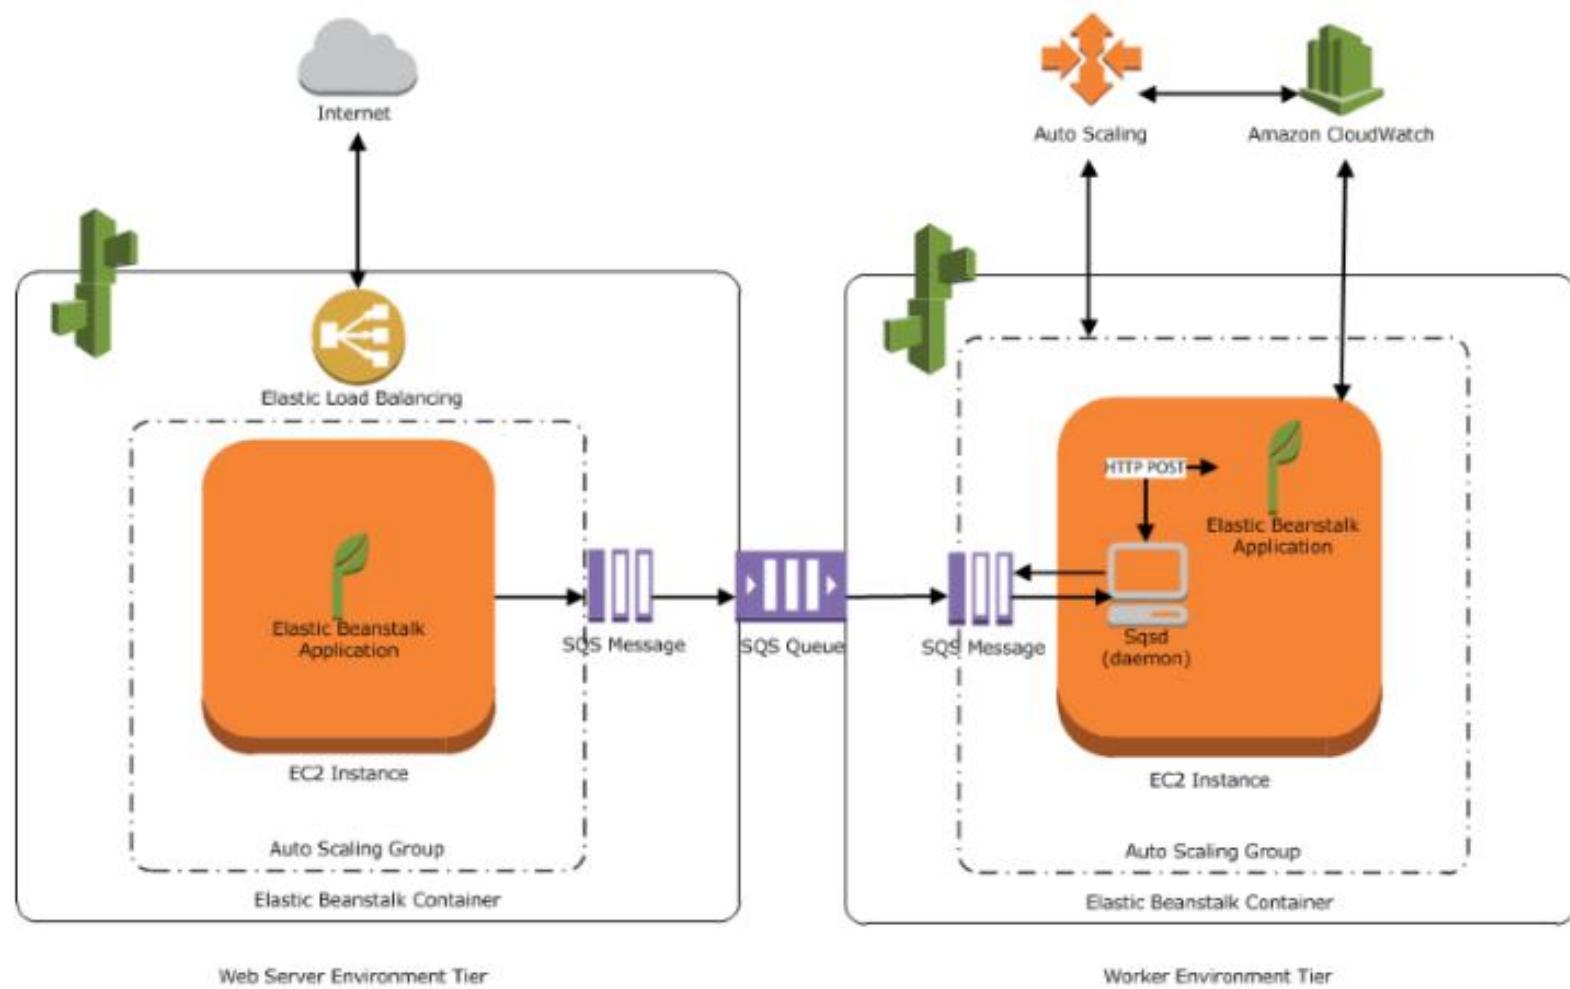
# Elastic Beanstalk Workflow



# Elastic Beanstalk Architecture (Web)



# Elastic Beanstalk Architecture (Worker)



# Elastic Beanstalk Competitors

- Cloud Foundry
- Google App Engine
- Azure Web Sites
- Heroku
- Bluemix

# Elastic Load Balancer

# Elastic Load Balancer



Elastic Load  
Balancing

- **Distributes** traffic across multiple EC2 instances, in multiple Availability Zones
- Supports **health checks** to detect unhealthy Amazon EC2 instances
- Supports the **routing and load balancing** of HTTP, HTTPS, SSL, and TCP traffic to Amazon EC2 instances

# Elastic Load Balancer Types

**Classic Load Balancer**

**PREVIOUS GENERATION**  
for HTTP, HTTPS, and TCP

Choose a Classic Load Balancer when you have an existing application running in the EC2-Classic network.

**Application Load Balancer**



HTTP  
HTTPS

Choose an Application Load Balancer when you need a flexible feature set for your web applications with HTTP and HTTPS traffic. Operating at the request level, Application Load Balancers provide advanced routing and visibility features targeted at application architectures, including microservices and containers.

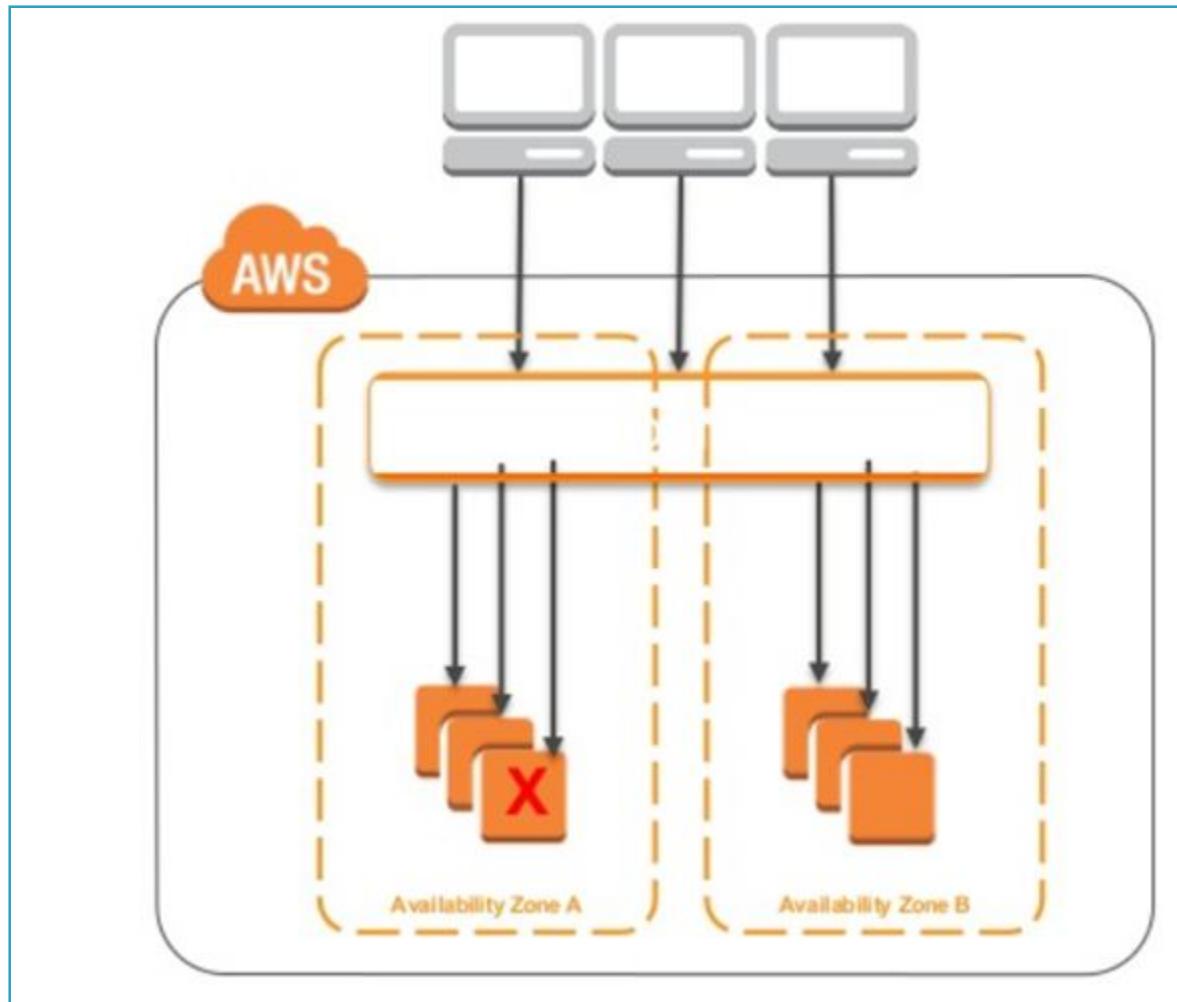
**Network Load Balancer**



TCP  
TLS

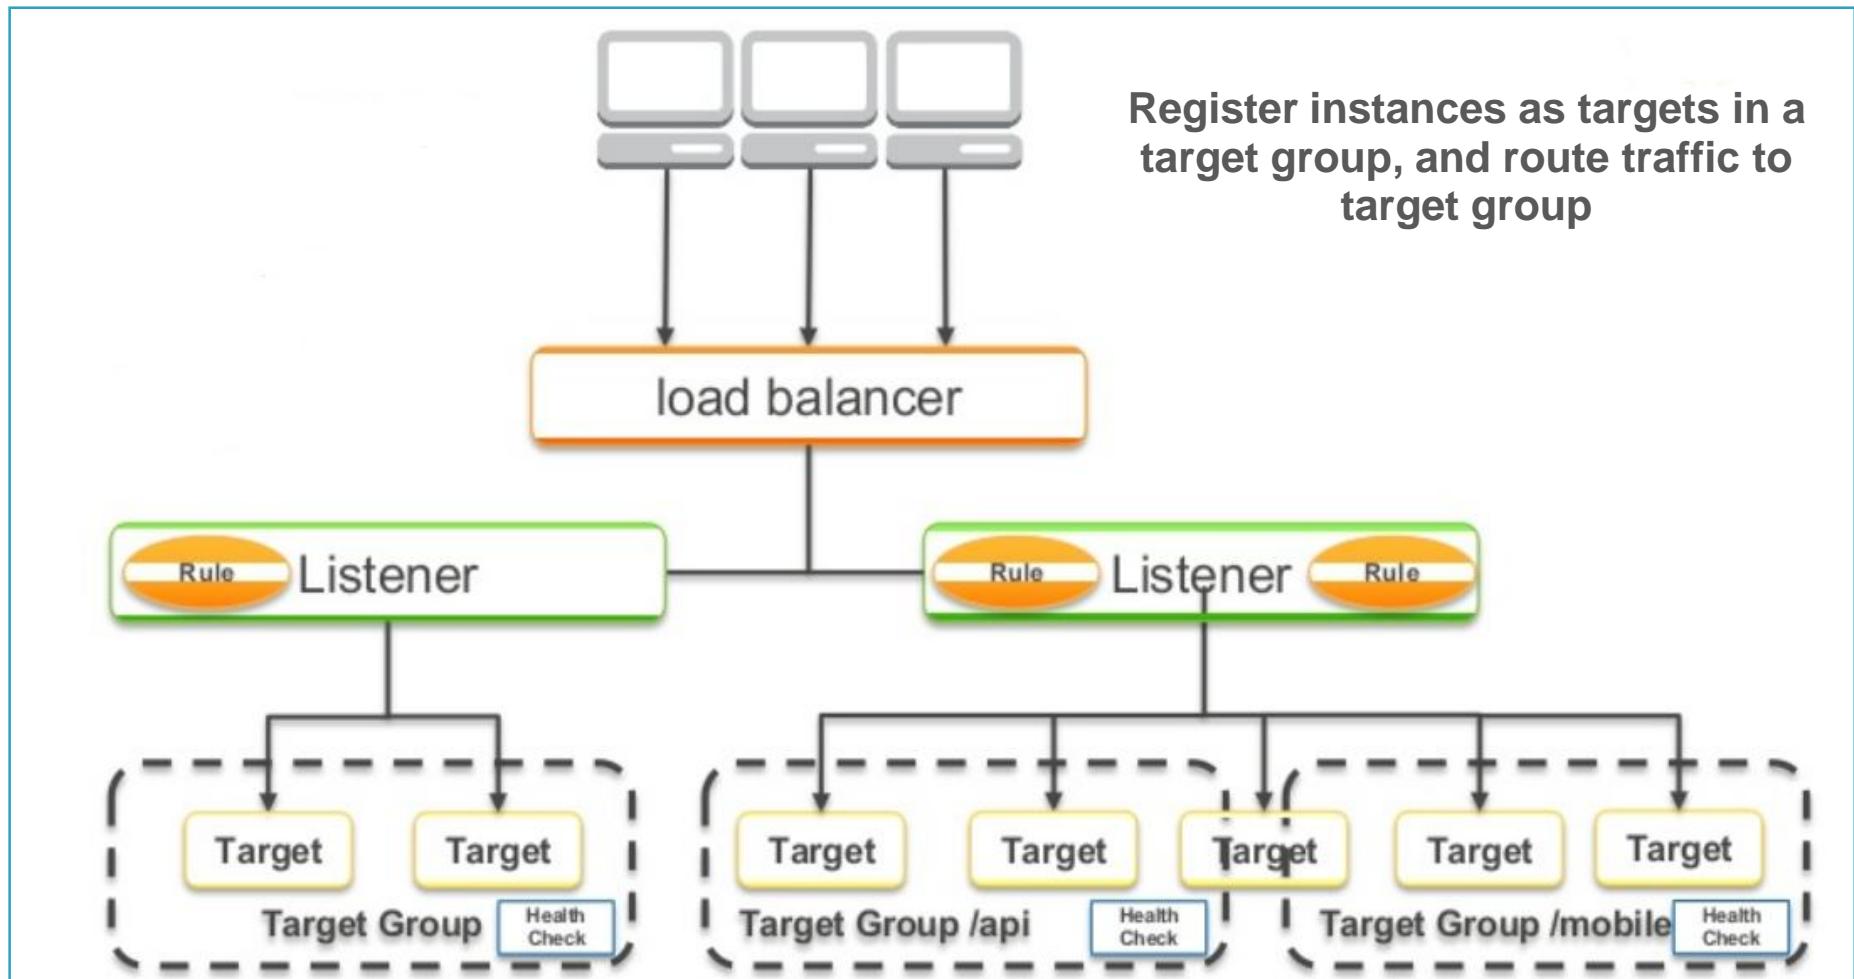
Choose a Network Load Balancer when you need ultra-high performance, the ability to terminate TLS connections at scale, centralize certificate deployment, and static IP addresses for your application. Operating at the connection level, Network Load Balancers are capable of handling millions of requests per second securely while maintaining ultra-low latencies.

# Classic Load Balancer



Register instance with your load balancer

# Application Load Balancer



# Classic vs Application Load Balancer

## Classic Load Balancer

benefits include support for:

- EC2-Classic.
- VPC.
- TCP and SSL listeners.
- Sticky sessions.
- OSI Layer 4  
*(network protocol level)*

## ALB benefits include support

for:

- Path-based routing.
- Routing requests to multiple services on a single EC2 instance.
- Containerized applications.
- Monitoring the health of each service independently.
- OSI Layer 7  
*(application level)*

# Auto Scaling

# Auto Scaling



Auto  
Scaling

- **Scale your Amazon EC2 capacity automatically**
- Well-suited for applications that experience **variability in usage**
- Available at no additional charge

# Auto Scaling Benefits

Better Fault Tolerance



Better Availability

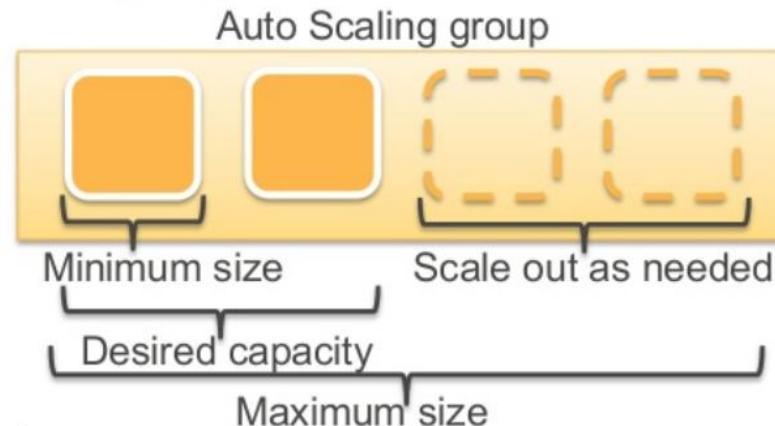


Better Cost Management



# Auto Scaling Groups

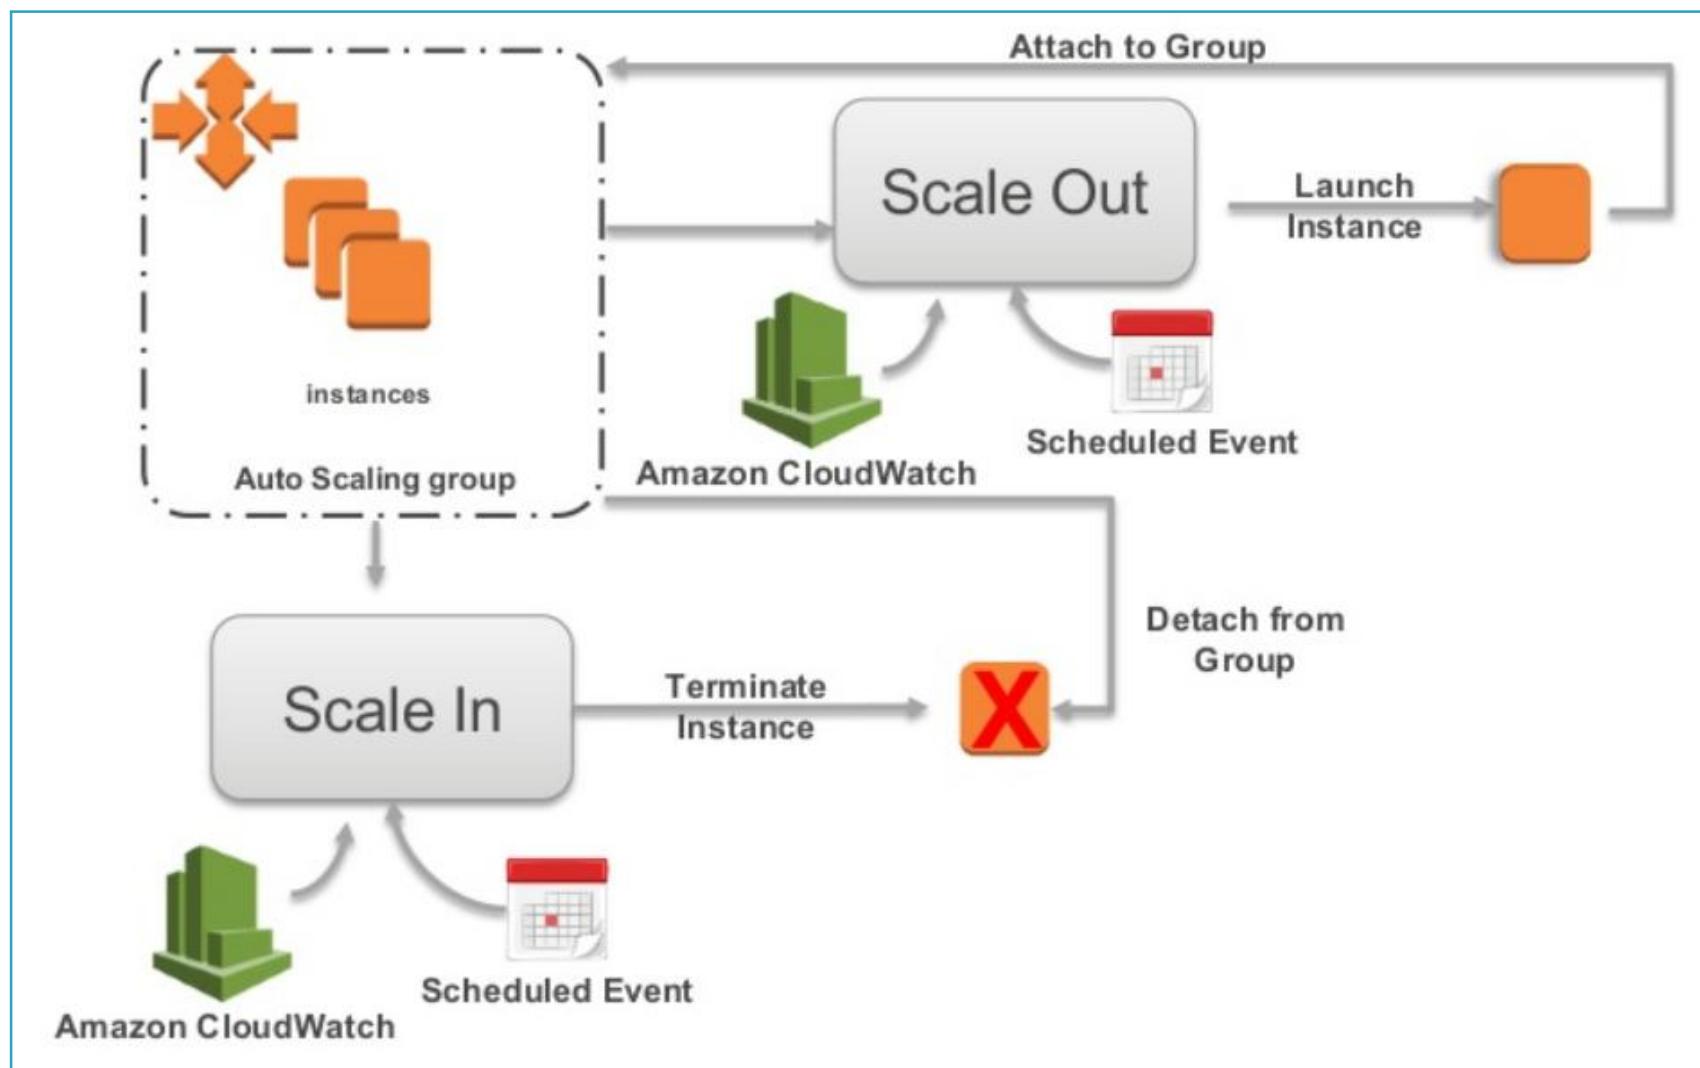
- Contain a collection of EC2 instances that share similar characteristics.
- Instances in an Auto Scaling group are treated as a **logical grouping** for the purpose of instance scaling and management.



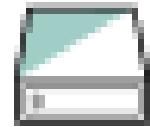
# Dynamic Scaling

- You can create a scaling policy that uses **CloudWatch alarms** to determine:
  - When your Auto Scaling group should **scale out**.
  - When your Auto Scaling group should **scale in**.
- You can use alarms to monitor:
  - Any of the metrics that AWS services send to Amazon CloudWatch.
  - Your own **custom metrics**.

# Auto Scaling Lifecycle



# AWS Storage Services



## Storage

Amazon Simple Storage Service (S3)

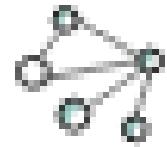
Amazon Elastic Block Storage (EBS)

Amazon Elastic File System (EFS)

Amazon Glacier

AWS Storage Gateway

# AWS Network Services



## Networking & Content Delivery

VPC

CloudFront

Direct Connect

Route 53

# AWS Database Services



# Database

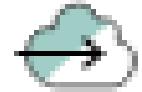
RDS

DynamoDB

ElastiCache

Amazon Redshift

# AWS Migration Services



## Migration

AWS Migration Hub

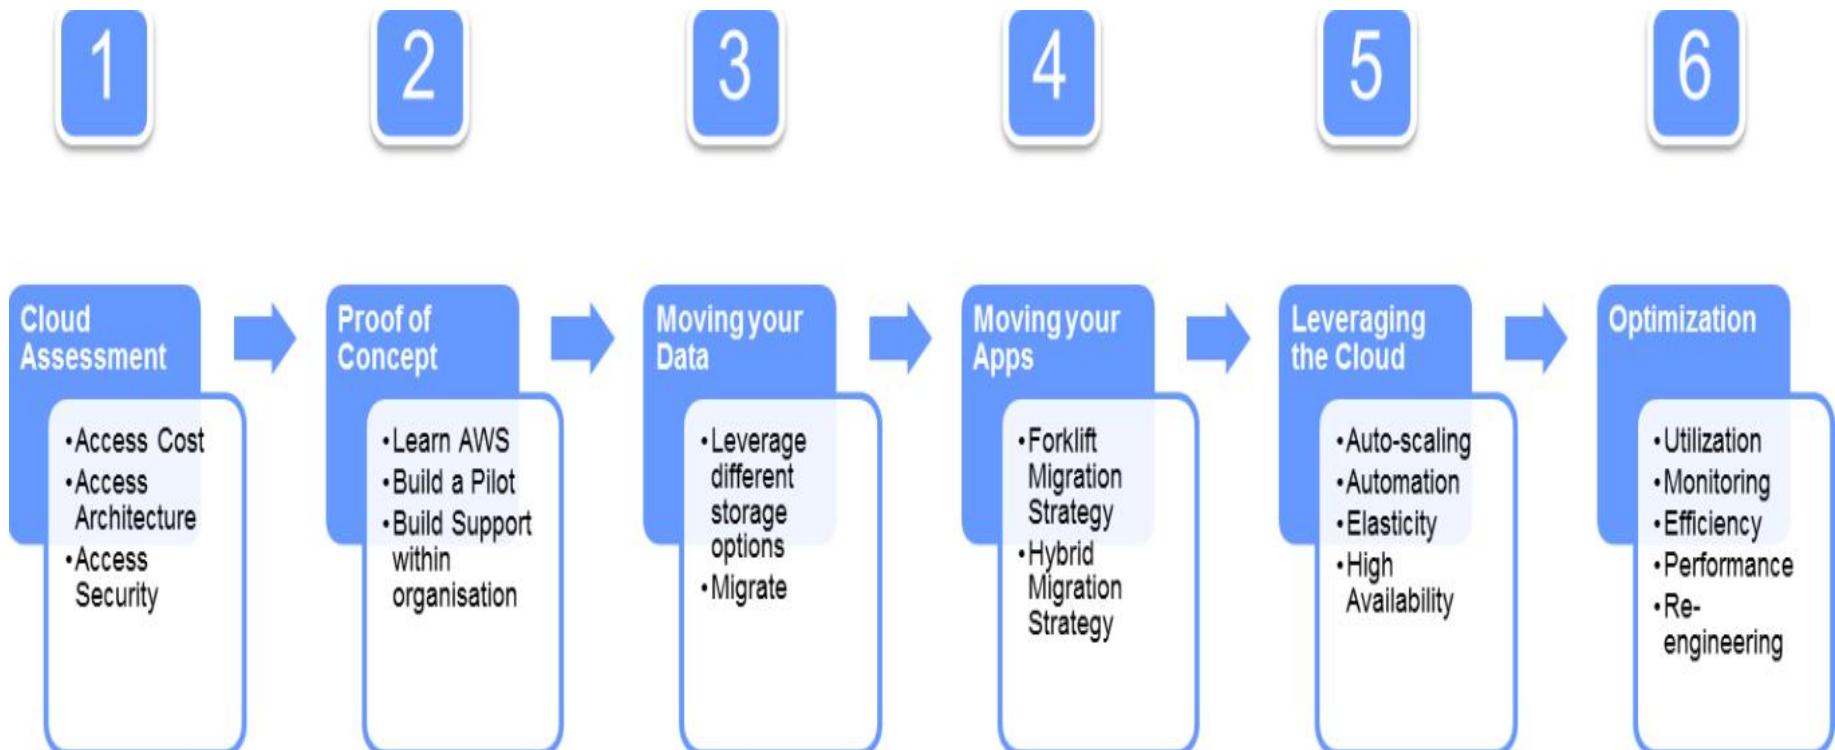
Application Discovery Service

Database Migration Service

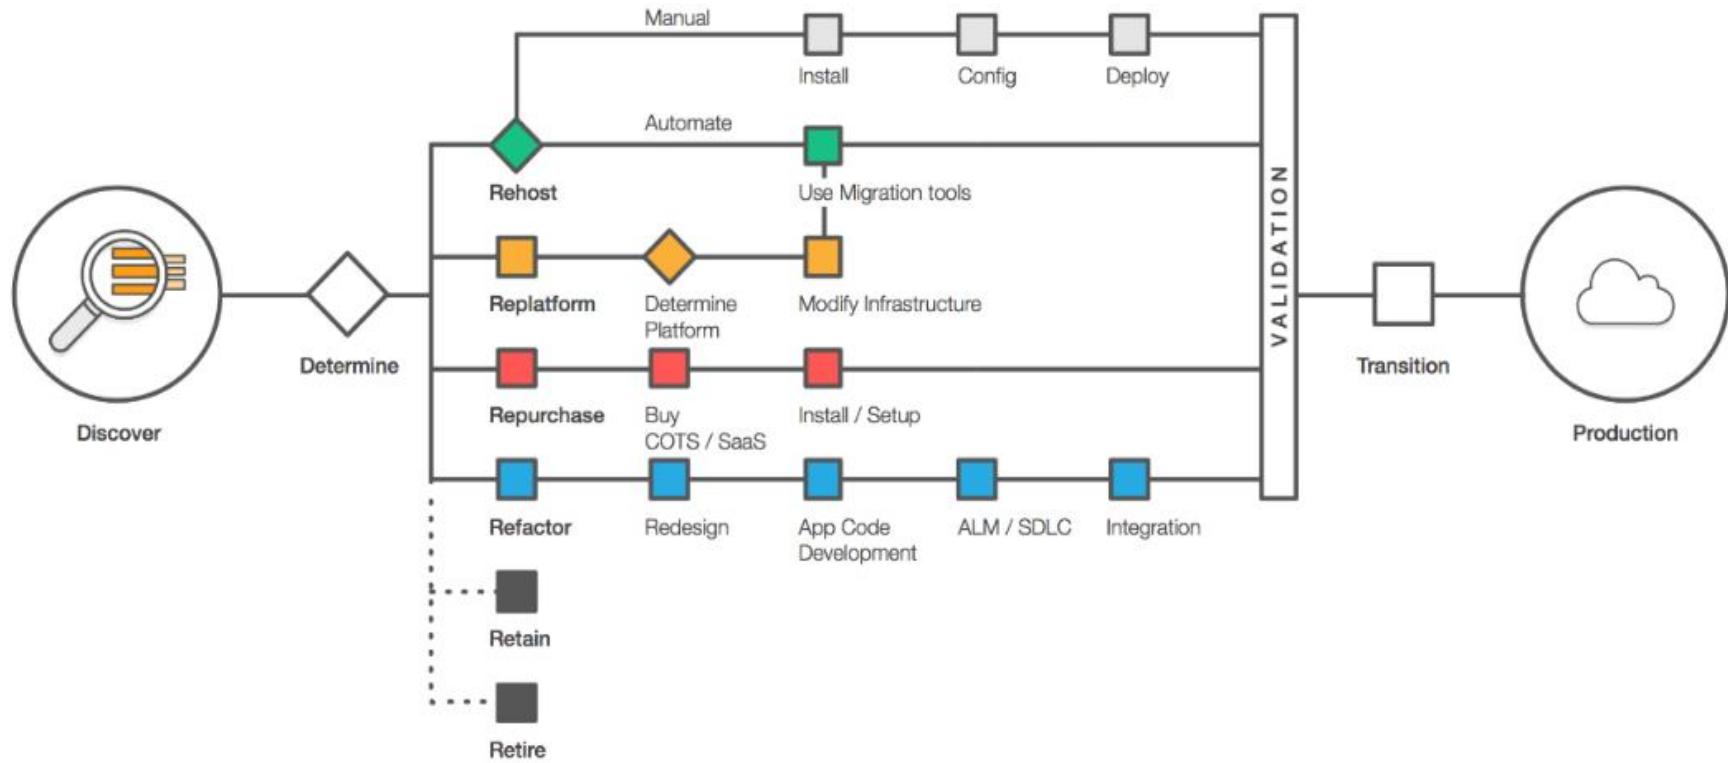
Server Migration Service

Snowball

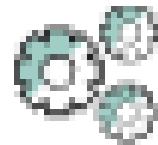
# Cloud Migration Strategy



# Cloud Migration Strategies



# AWS Developer Services



## Developer Tools

CodeStar

CodeCommit

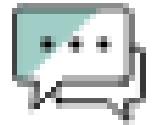
CodeBuild

CodeDeploy

CodePipeline

X-Ray

# AWS Messaging Services



## Messaging

Simple Queue Service

Simple Notification Service

Simple Email Service

# AWS Security Services



## Security, Identity & Compliance

IAM

Inspector

Certificate Manager

Directory Service

WAF & Shield

Artifact

Amazon Macie

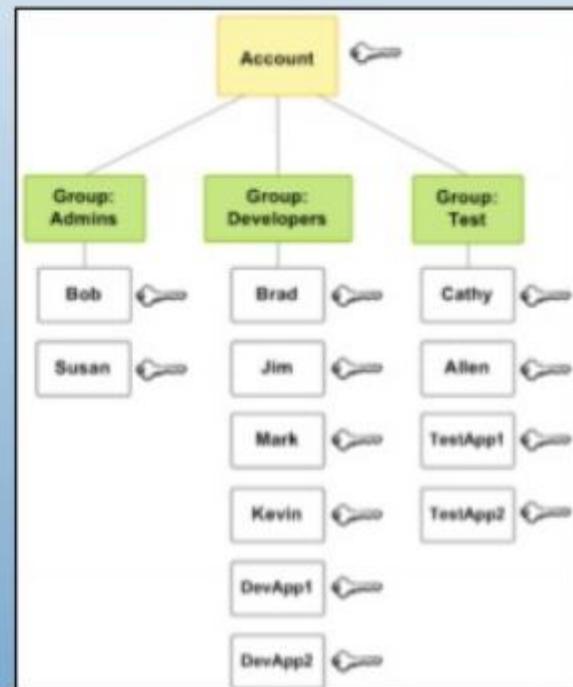
CloudHSM

# AWS Identity and Access Management (IAM)

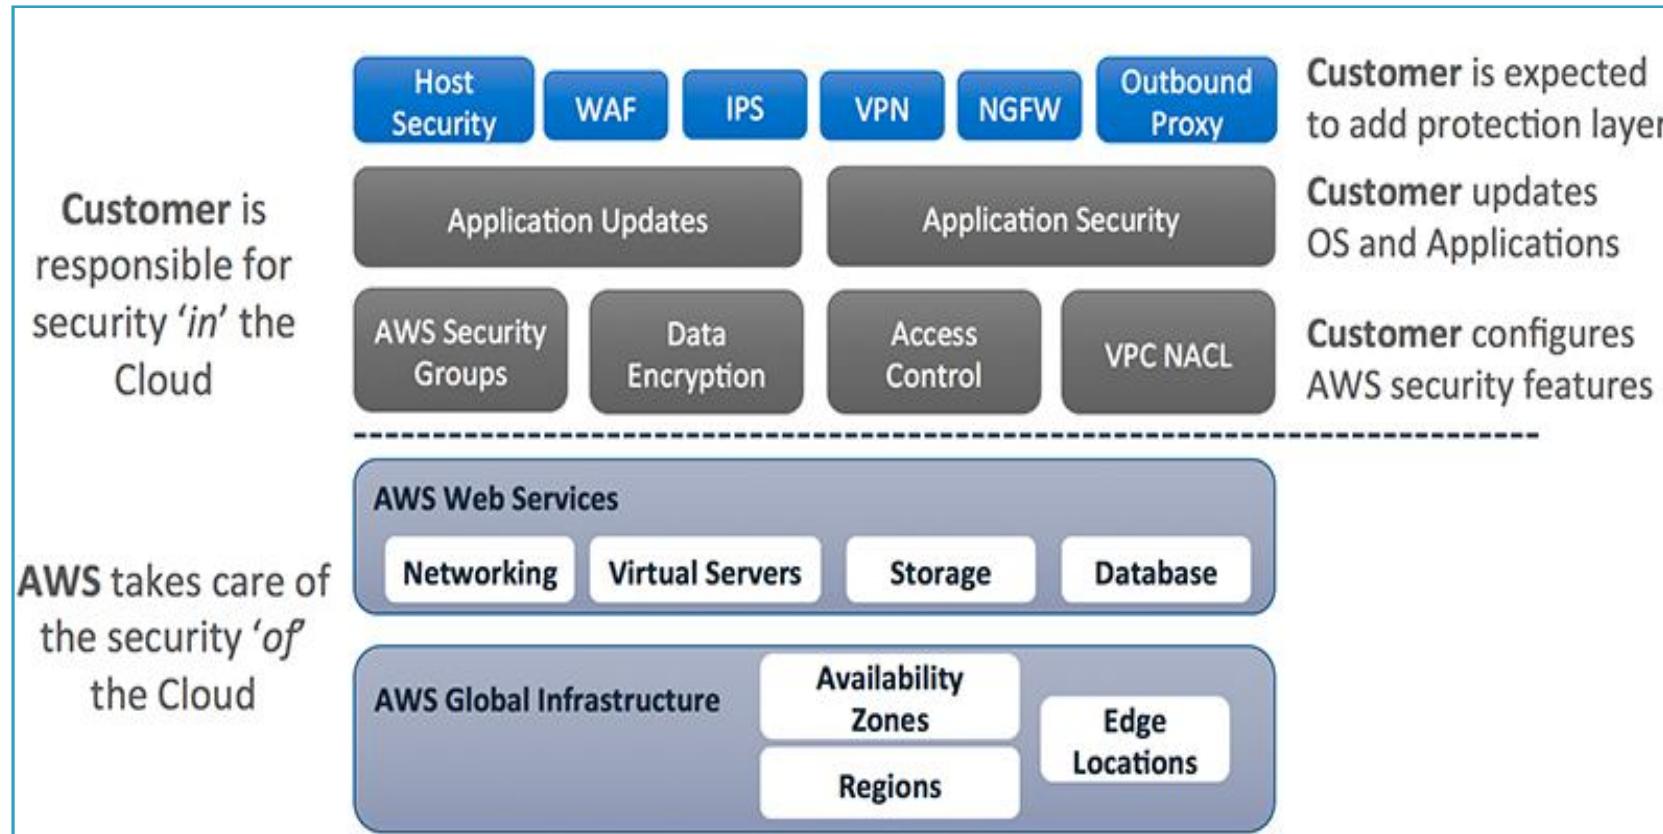
**IAM** is a web service for securely controlling access to AWS services. Centrally manage users, security credentials such as access keys, and permissions that control which AWS resources users and applications can access.

# AWS Identity and Access Management (IAM)

- Each account has root identity plus Users, Groups, Roles
  - Account-level: password complexity policies
- Unique security credentials for each user
  - Login/password (optional)
  - Access / secret keys (for APIs) (optional)
  - (V)MFA devices (optional)
- Policies control access to AWS APIs
- Deeper integration into some Services
  - S3: policies on objects and buckets
  - Simple DB: domains
- AWS Management Console supports IAM user log on
- Not for Operating Systems or Applications
  - use LDAP, Active Directory/ADFS, etc...



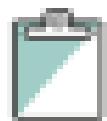
# AWS Shared Security Model



# AWS Security Compliance

| Certificates:   | Programmes:               |
|---|---------------------------|
|    | PCI DSS Level 1           |
|    | SOC 1 / ISAE 3402         |
|    | SOC 2                     |
|    | SOC 3                     |
|    | ISO 9001                  |
|   | IRAP (Australia)          |
|   | FIPS 140-2                |
|   | CJIS                      |
|  | CSA                       |
|  | FERPA                     |
|  | HIPAA                     |
|   | FedRAMP (SM)              |
|   | DoD CSM Levels 1-2, 3-5   |
|   | DIACAP and FISMA          |
|   | ISO 27001                 |
|   | MTCS Tier 3 Certification |
|   | ITAR                      |
|   | MPAA                      |
|   | G-Cloud                   |
|   | Section 508 / VPAT        |

# AWS Management Services



## Management Tools

CloudWatch

CloudFormation

CloudTrail

Config

OpsWorks

Service Catalog

Trusted Advisor

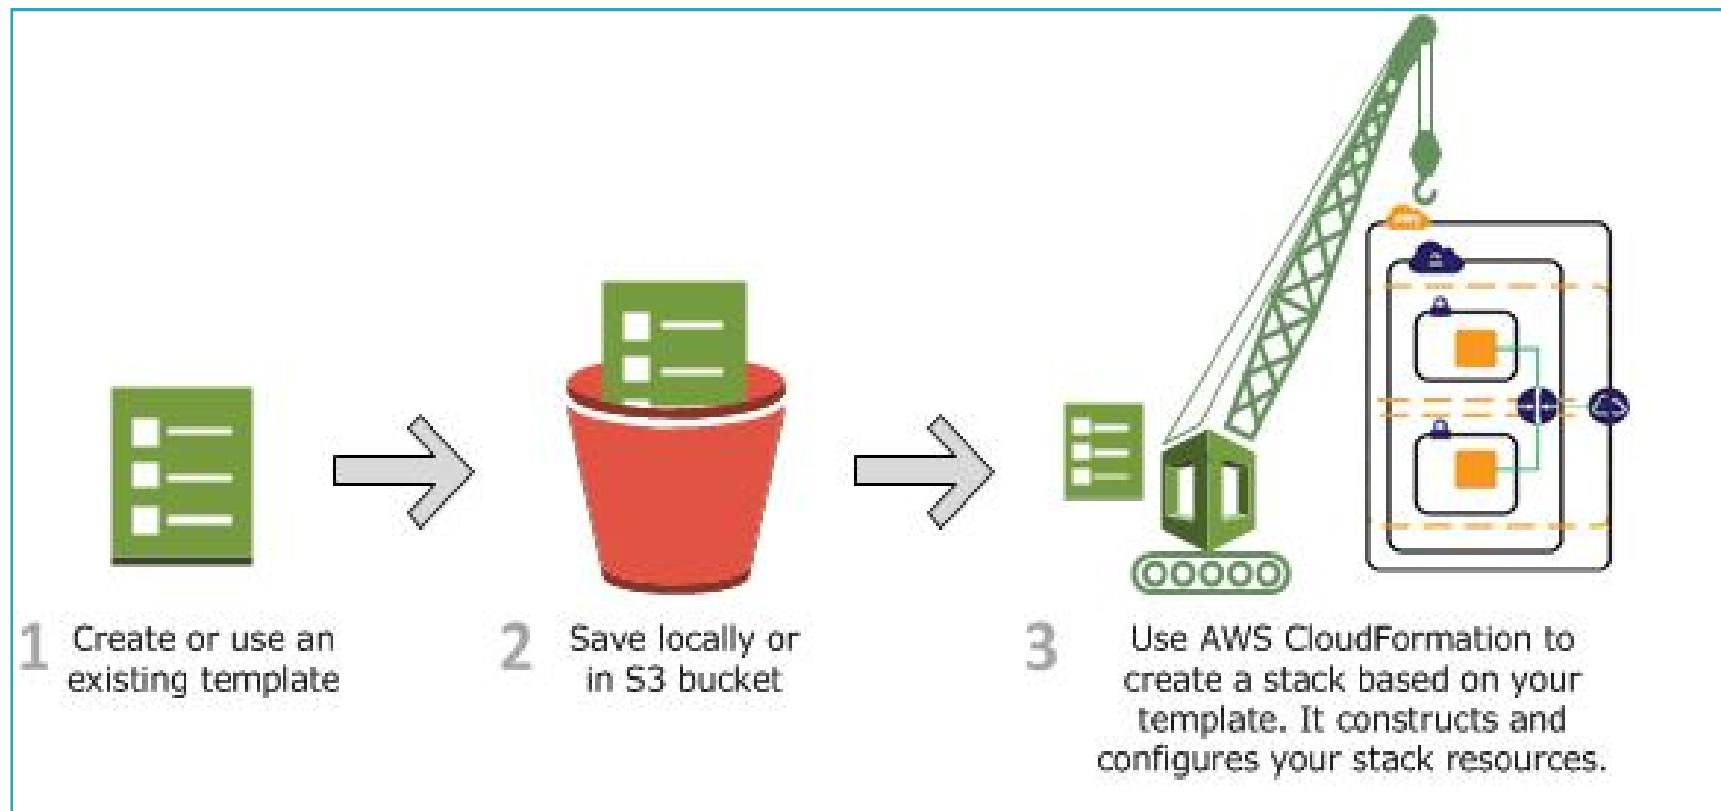
Managed Services

# Cloud Formation

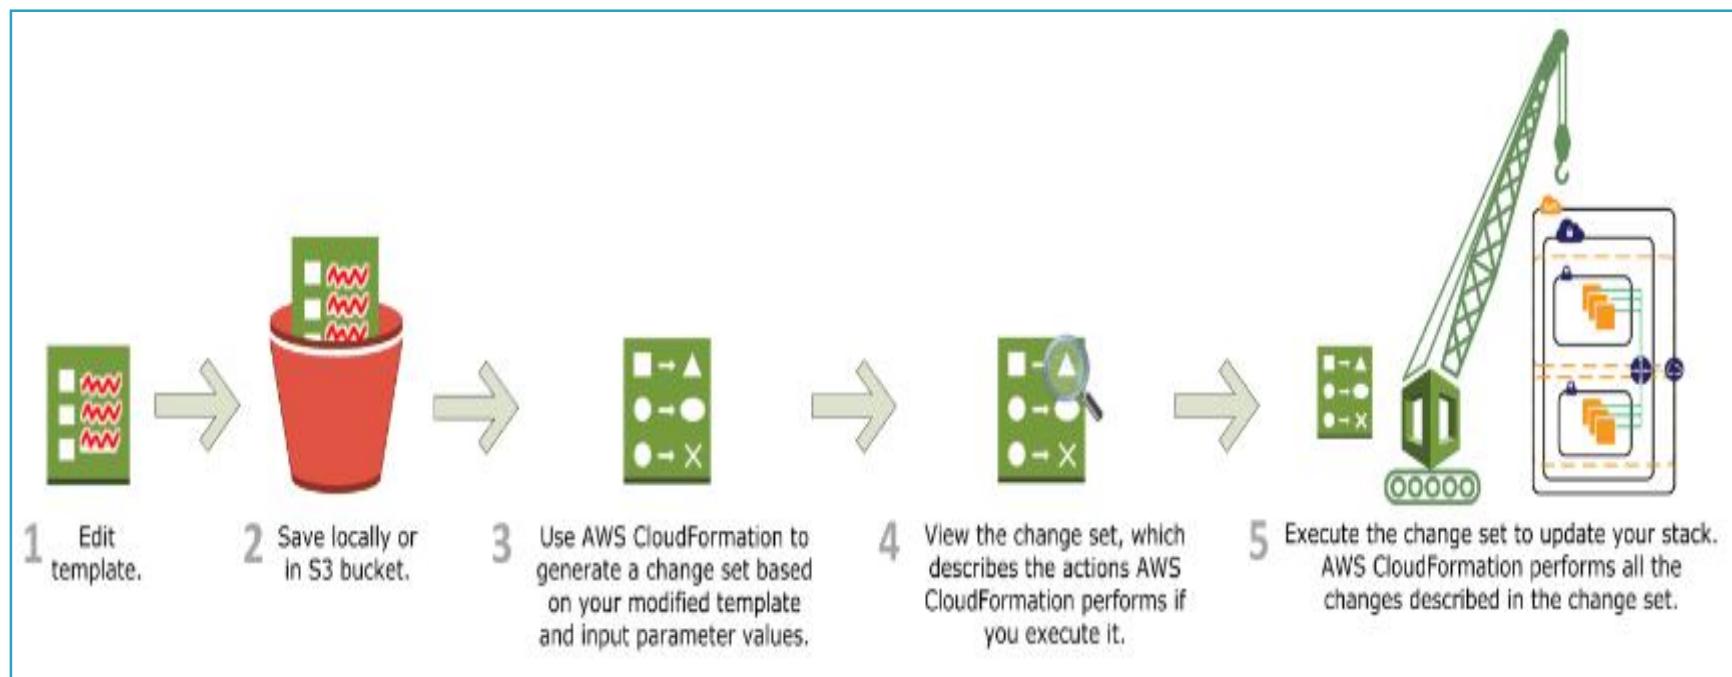
AWS CloudFormation enables you to create and provision AWS infrastructure deployments predictably and repeatedly

- **Templates** – Blueprints for building AWS resources (JSON or YAML)
- **Stack** – Logical group of related resources to be managed as single unit
- **Change Sets** – Shows how changes might impact the running resources

# Cloud Formation – Create Process



# Cloud Formation – Update Process



# Cloud Formation – Best Practices

## Planning and organizing

- Organize Your Stacks By Lifecycle and Ownership
- Use Cross-Stack References to Export Shared Resources
- Use IAM to Control Access
- Reuse Templates to Replicate Stacks in Multiple Environments
- Verify Quotas for All Resource Types
- Use Nested Stacks to Reuse Common Template Patterns

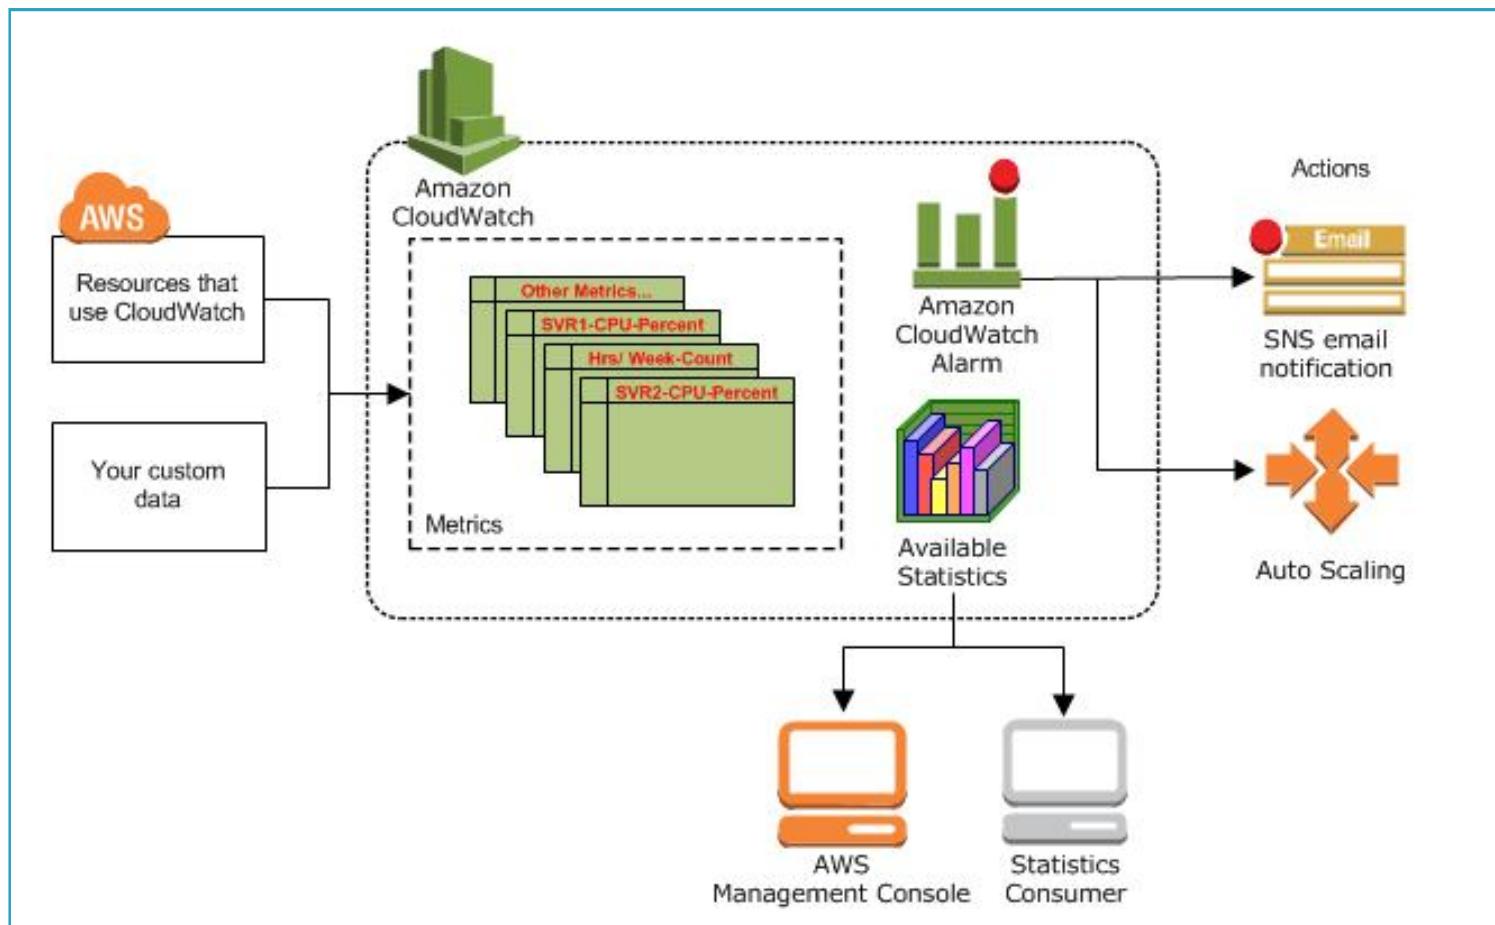
## Creating templates

- Do Not Embed Credentials in Your Templates
- Use AWS-Specific Parameter Types
- Use Parameter Constraints
- Use AWS::CloudFormation::Init to Deploy Software Applications on Amazon EC2 Instances
- Use the Latest Helper Scripts
- Validate Templates Before Using Them

## Managing stacks

- Manage All Stack Resources Through AWS CloudFormation
- Create Change Sets Before Updating Your Stacks
- Use Stack Policies
- Use AWS CloudTrail to Log AWS CloudFormation Calls
- Use Code Reviews and Revision Controls to Manage Your Templates
- Update Your Amazon EC2 Linux Instances Regularly

# Cloud Watch



# AWS Other Services



## Analytics

Athena  
EMR  
CloudSearch  
Elasticsearch Service  
Kinesis  
Data Pipeline  
QuickSight  
AWS Glue



## Artificial Intelligence

Lex  
Amazon Polly  
Rekognition  
Machine Learning



## Application Services

Step Functions  
SWF  
API Gateway  
Elastic Transcoder



## Internet Of Things

AWS IoT  
AWS Greengrass



## Business Productivity

WorkDocs  
WorkMail  
Amazon Chime



## Mobile Services

Mobile Hub  
Cognito  
Device Farm  
Mobile Analytics  
Pinpoint



## Game Development

Amazon GameLift



## Contact Center

Amazon Connect

# Thank You!